

## Prepare dataset [1]


I created a new file `prepare_data_mpi_inf_3dhp.py`. It is based on `prepare_data_h36m.py`. I basically structured the file into the same format as described above with one important difference. For Human3.6M, the ground truth 3D joints are recorded with World coordinate. Thus, during the data processing step. The 3D joint coordinates need to be converted into Camera Coordinates. However, `Mpi_inf_3dhp` is recorded in Camera coordinates already, so I do not need to convert them. Thus, I made some changes to the training script to reflect this.

```
54 print('Preparing data...')
55 if args.dataset.startswith('mpi'):
56     for subject in dataset.subjects():
57         for action in dataset[subject].keys():
58             anim = dataset[subject][action]
59
60             if 'positions' in anim:
61                 anim['positions_3d'] = anim['positions']
62 else:
63     for subject in dataset.subjects():
64         for action in dataset[subject].keys():
65             anim = dataset[subject][action]
66
67             if 'positions' in anim:
68                 positions_3d = []
69                 for cam in anim['cameras']:
70                     pos_3d = world_to_camera(anim['positions'], R=cam['orientation'], t=cam['translation'])
71                     pos_3d[:, 1:] -= pos_3d[:, :1] # Remove global offset, but keep trajectory in first position
72                     positions_3d.append(pos_3d)
73                 anim['positions_3d'] = positions_3d
74
```

## Dataset class

Next, I mimic the structure of `h36m_dataset.py` to make `mpi_inf_3dhp_dataset.py`. Three important changes I need to make are fps, skeleton, and camera. For fps, the paper [] documents the fps for `MPI_INF_3DHP` as 25fps.

For skeleton, `MPI_INF_3DHP` provides a script that can convert the index to Human3.6M style index. Then, I am able to reuse the human3.6M skeleton model since it has the matching joint index. I also made a drawing to show the index.



Lastly, I need to handle the camera parameters.

To project 3D joints to 2D. It requires camera intrinsic parameters. For visualization purposes, it also needs extrinsic parameters. When looking through the h36m dataset. I found out the following intrinsic parameters are needed.

```
h36m_cameras_intrinsic_params = [
    {
        'id': '54138969',
        'center': [512.54150390625, 515.4514770507812],
        'focal_length': [1145.0494384765625, 1143.7811279296875],
        'radial_distortion': [-0.20709891617298126, 0.24777518212795258, -0.0030751503072679043],
        'tangential_distortion': [-0.0009756988729350269, -0.00142447161488235],
        'res_w': 1000,
        'res_h': 1002,
        'azimuth': 70, # Only used for visualization
    },
],
```

I need a center, focal length, radial and tangential distortion, width and height.

For extrinsic parameters, I need to specify the orientation and the translation.

```
h36m_cameras_extrinsic_params = {
    's1': [
        {
            'orientation': [0.1407056450843811, -0.1500701755285263, -0.755240797996521, 0.6223280429840088],
            'translation': [1841.1070556640625, 4955.28466796875, 1563.4454345703125],
        },
        {
            'orientation': [0.6157187819480896, -0.764836311340332, -0.14833825826644897, 0.11794740706682205],
            'translation': [1761.278564453125, -5078.0068359375, 1606.2650146484375],
        },
        {
            'orientation': [0.14651472866535187, -0.14647851884365082, 0.7653023600578308, -0.6094175577163696],
            'translation': [-1846.7777099609375, 5215.04638671875, 1491.972412109375],
        },
        {
            'orientation': [0.5834008455276489, -0.7853162288665771, 0.14548823237419128, -0.14749594032764435],
            'translation': [-1794.7896728515625, -3722.698974609375, 1574.8927001953125],
        },
    ],
},
```

Within the MPI\_INF\_3DHP, I've been given the following parameters.

```
2 name      0
3 sensor    10 10
4 size      2048 2048
5 animated   0
6 intrinsic 1497.693 0 1024.704 0 0 1497.103 1051.394 0 0 0 1 0 0 0 0 1
7 extrinsic 0.9650164 0.00488022 0.262144 -562.8666 -0.004488356 -0.9993728 0.0351275 1398.138 0.262151 -0.03507521 -0.9643893 3852.623 0 0 0 1
8 radial     0
```

We have size, intrinsic matrix and extrinsic matrix.

The size specifies the width and height.

First of all, let's take a look at the extrinsic parameters. I am putting the line into matrix form.

0.9650164	0.00488022	0.262144	-562.8666
-0.004488356	-0.9993728	0.0351275	1398.138
0.262151	-0.03507521	-0.9643893	3852.623
0	0	0	1

In this website [2], it talks about the camera extrinsic:

The extrinsic matrix takes the form of a rigid transformation matrix: a 3x3 rotation matrix in the left-block, and 3x1 translation column-vector in the right:

$$[R | t] = \left[ \begin{array}{ccc|c} r_{1,1} & r_{1,2} & r_{1,3} & t_1 \\ r_{2,1} & r_{2,2} & r_{2,3} & t_2 \\ r_{3,1} & r_{3,2} & r_{3,3} & t_3 \end{array} \right]$$

It's common to see a version of this matrix with extra row of (0,0,0,1) added to the bottom. This makes the matrix square, which allows us to further decompose this matrix into a rotation *followed* by translation:

$$\begin{aligned} \left[ \begin{array}{ccc|c} R & t \\ \hline 0 & 0 & 0 & 1 \end{array} \right] &= \left[ \begin{array}{ccc|c} I & t \\ \hline 0 & 0 & 0 & 1 \end{array} \right] \times \left[ \begin{array}{ccc|c} R & 0 \\ \hline 0 & 0 & 0 & 1 \end{array} \right] \\ &= \left[ \begin{array}{ccc|c} 1 & 0 & 0 & t_1 \\ 0 & 1 & 0 & t_2 \\ 0 & 0 & 1 & t_3 \\ \hline 0 & 0 & 0 & 1 \end{array} \right] \times \left[ \begin{array}{ccc|c} r_{1,1} & r_{1,2} & r_{1,3} & 0 \\ r_{2,1} & r_{2,2} & r_{2,3} & 0 \\ r_{3,1} & r_{3,2} & r_{3,3} & 0 \\ \hline 0 & 0 & 0 & 1 \end{array} \right] \end{aligned}$$

Thus, we can see the upper 3\*3 is the rotation matrix, while the three variables at right are the translational matrix.

However, this is not the format I wanted. Rather than a 3x3 matrix, I need a Quaternion.

Thanks to this website [3], I am able to make the conversion.

### 3D Rotation Converter

#### Input

Input angle format ☒ Radians ☐ Degrees

Rotation matrix		
<input type="text" value="0.9650164"/>	<input type="text" value="0.00488022"/>	<input type="text" value="0.262144"/>
<input type="text" value="0.04488356"/>	<input type="text" value="-0.999372"/>	<input type="text" value="0.0351275"/>
<input type="text" value="0.262151"/>	<input type="text" value="0.03507521"/>	<input type="text" value="0.9643893"/>

Quaternion

x  y  z  w (real part)

Axis-angle

Axis x  y  z  Angle (radians)

Axis with angle magnitude (radians)

Axis x  y  z

Euler angles of multiple axis rotations (radians)

XYZ  x  y  z

Triple of points, P, Q, R, such that X  $\parallel$  (Q-P), Z  $\parallel$  X  $\times$  (R-P), and Y  $\parallel$  Z  $\times$  X.

P: x  y  z

Q: x  y  z

R: x  y  z

#### Output

Output angle format ☒ Radians ☐ Degrees

Rotation matrix

```
[ 0.9650164, 0.0048802, 0.2621440;
-0.04488356, -0.9993728, 0.0351275;
0.2621510, -0.0350752, 0.9643892 ]
```

Quaternion [x, y, z, w]

```
[ 0.9910573, 0.0000989, 0.1322565, -0.017709 ]
```

Axis-Angle [[x, y, z], angle (radians)]

```
{ [ 0.9912128, 0.0000989, 0.1322772 ], 3.1770126 }
```

Axis with angle magnitude (radians) [x, y, z]

```
[ 3.1490954, 0.0003141, 0.4202464 ]
```

Euler angles (radians) XYZ

```
[ x: -3.1051841, y: 0.2652432, z: -0.0050571 ]
```

This handles the extrinsic parameters.

As for the intrinsic parameters, I will first put the line into matrix form.

1497.693	0	1024.704	0
0	1497.103	1051.394	0
0	0	1	0
0	0	0	1

This website[4] talks about the intrinsics.


$$K = \begin{pmatrix} f_x & s & x_0 \\ 0 & f_y & y_0 \\ 0 & 0 & 1 \end{pmatrix}$$

$f_x$  and  $f_y$  are the focal length,  $x_0$  and  $y_0$  are the center.

The last parameter I need are the radial and tangential distortion.

Radial distortion is specified as 0 in the calibration file. However, it did not specify the tangential distortion, I will assume it is 0. The resulting code is available on my github page:

## Result and Discussion



The resulting MPJPE for MPI\_INF\_3DHP is at around 80mm. The figure above shows the training curve. It does have some spikes which is due to the improper momentum value set for the batch normalization layer. MPI\_INF\_3DHP is smaller than HumanPose3.6M, thus it needs a higher momentum value which also needs a smaller architecture. The training parameter I set for the above graph is default 3,3,3 arc which is 27 frames. Within the run code, I modified the momentum to be around 0.8 to 0.6. The resulting graph is not that perfect, but it can be fixed with some parameter tuning. Thus, my work shows that VideoPose3D can also be generalized on MPI\_INF\_3DHP dataset.

## References

1. Pavllo, D., Feichtenhofer, C., Grangier, D., & Auli, M. (2019). 3D human pose estimation in video with temporal convolutions and semi-supervised training. In Conference on Computer Vision and Pattern Recognition (CVPR).
2. Simek, K. (2012, August 22). *Dissecting the camera matrix, part 2: The extrinsic matrix*. Retrieved April 29, 2022, from <https://ksimek.github.io/2012/08/22/extrinsic/>
3. Gaschler, A. (n.d.). 3D Rotation Converter. Retrieved April 29, 2022, from <https://www.andre-gaschler.com/rotationconverter/>
4. Simek, K. (2013, August 13). *Dissecting the camera matrix, part 3: The Intrinsic matrix*. Retrieved April 29, 2022, from <https://ksimek.github.io/2013/08/13/intrinsic/>